

SWAPNIL MITTAL

swapnil.mittal1@gmail.com | linkedin.com/in/mittalswapnil | github.com/swapnilmittal1 | swapnilmittal.com | 404-513-8227

EDUCATION

Georgia Institute of Technology

Master of Science in Computer Science, Concentration: Machine Learning

Atlanta, Georgia

Expected December 2026

Georgia Institute of Technology

Bachelor of Science in Computer Science (Dean's List)

Atlanta, Georgia

December 2025

TECHNICAL SKILLS

Programming Languages: Python, Java, C++, Go, Rust, TypeScript, JavaScript, SQL, Bash

Frameworks & Libraries: PyTorch, TensorFlow, Transformers (Hugging Face); React, Node.js, Express.js, Flask, Django

Cloud & DevOps: AWS, Azure, GCP; Docker, Kubernetes, Terraform, Ansible, Helm; GitHub Actions, Jenkins

Databases & Tools: PostgreSQL, MySQL, MongoDB, Redis, SQLite; Git; Linux; CI/CD; REST APIs, GraphQL

EXPERIENCE

Orbex Labs

Machine Learning Engineer Intern

Remote

May 2025 – Present

- Built end-to-end document processing pipeline (OCR of 100K+ scanned briefs, spaCy normalization, centralized feature store) and fine-tuned RoBERTa-Legal for clause extraction/classification, achieving 92% F1 on indemnity.
- Developed a RAG-based Q&A assistant integrated into the client portal using LangChain and Cohere embeddings and implemented RESTful APIs and CI/CD pipelines, reducing attorney query resolution time by 40%.
- Created React human-in-the-loop annotation tools with unit testing in an Agile framework and prototyped generative AI pipelines with LLaMA models on a legal knowledge graph, raising NER accuracy from 65% to 95.2%.

Georgia Tech — Create-X Capstone

Head Teaching Assistant

Atlanta, GA

January 2025 – May 2025

- Coordinate a cohort of 100 interdisciplinary student engineers and 5–10 faculty mentors, and managing course logistics to ensure smooth delivery.
- Answer mentor questions, guide student teams through problem validation, product design iterations, MVP rollouts, and potential commercialization pathways.

KBR, Inc.

Software Engineering Intern

Houston, TX

May 2024 – Aug 2024

- Engineered LLM microservices on Azure AKS (TensorFlow, PyTorch) with Docker, Helm & auto-scaling—500+ req/sec at 99.999% uptime; set up CI/CD (GitHub Actions) and Prometheus/Grafana monitoring.
- Built 2 TB/day ETL pipeline (Azure Data Lake, Synapse, Data Factory) with Terraform IaC & GitOps in Agile/Scrum—ensured 99.99% fidelity, real-time anomaly detection & quality checks.
- Developed HPC simulations (Python, VIKTOR, MATLAB, AutoCAD, RFEM)—boosted computational precision 40%, reduced runtimes 35%, and improved cross-team collaboration via RESTful APIs, and peer code reviews.
- Optimized ETL workflows (parallel processing, intelligent caching, Terraform IaC)—cut operational costs 12%, slashed pipeline latency 40%, and enhanced observability with centralized logging, and dashboard automation.

PROJECTS & OPEN-SOURCE

Reverse-Proxy API Debugger | C++, POCO, Palantir AIP, React, PostgreSQL

Jan 2025 – April 2025

- Wrote an Envoy-WASM sidcar in C++20 that mirrors 100% HTTP/gRPC traffic, decorates each request with OpenTelemetry spans, and streams traces to ClickHouse at 15k req/s with <200 ms p99 overhead.
- Built a RAG layer over OpenAPI & Confluence using LangChain and Cohere, improving precision to 97%.
- Developed a React/Grafana console for JSON insights, slashing MTTR from 3 h to 12 min, saving 20+ hrs/min.

AI Social Commerce Platform | Python, React, GraphQL, PyTorch, Pinecone, Meta APIs

Dec 2024 – May 2025

- Deployed federated GraphQL gateway (Go, Apollo) unifying Meta Graph, Messenger, and Shops on AWS Fargate.
- Fine-tuned LoRA-compressed LLaMA via LangChain & Cohere API—cut reply latency 47%
- Streamed Kafka→Flink→Druid events; built Next.js dashboard auto-budgeter, raising ROAS 25%.

Prompt+ | HackGT 11 Hackathon Overall Winner

Sep 2024

- Built a real-time prompt-optimization tool reducing LLM token usage by 47% via context analysis.
- Implemented dynamic UI in React + Tailwind and real-time collaboration with Socket.IO.